

# Written Assignment 3:

## Query Plan, and Query Optimization

Name: Venkata Lakshmi Sasank Tipparaju

Student# 700738838

CS 5600 – 13892

### 1. Join Algorithm (15 points)

Consider the join operation between relation  $r$  and  $s$  ( $r \bowtie_{\theta} s$ ),  $\theta$  is  $r.A = s.B$  with the following information:

Relation  $r$  contains 10,000 tuples and has 10 tuples per block.

Relation  $s$  contains 2,000 tuples and has 10 tuples per block.

There are 33 buffer blocks available in Memory.

No sorted data in relation  $r$  and  $s$ .

Find Total cost (block transfers)

- a. Using Block Nested Loop Join (3 points)
- b. Using Merge Join (2 points)
- c. Using Hash Join (Recursive partition) (2 points)
- d. Using Hash Join (No recursive partition) (3 points)
- e. If there are the infinity of memory, which join algorithm that you prefer? And why? (5 points)

## Written Assignment 3

Sasank

①

Given Relation b/w  $r$  &  $s$   
( $r \bowtie s$ ),  $\theta$  is  $r.A = s.B$

Relation ' $r$ ' contains 10,000 tuples & 10 tuples per block.

$$\Rightarrow \text{Total no. of blocks required for 'r'} = \frac{10000}{10} \\ = 1000 \text{ blocks}$$

Relation ' $s$ ' contains 2000 tuples & 10 tuples per block

$$\Rightarrow \text{Total no. of blocks required for 's'} = 2000/10 \\ = 200 \text{ blocks}$$

No. of buffer blocks available in memory = 33 blocks

$\therefore$   $r$  needs 1000 blocks,  $s$  needs 200 blocks,  $M=33$  blocks

(1.1)

Block Nested Loop Join:

$$\text{Cost} = \left( \left\lceil \frac{b_r}{M-1} \right\rceil \cdot b_s \right) + b_r \cdot \text{block\_transfer}$$

$$b_r = 1000, \quad b_s = 200, \quad m = 33$$

$$\Rightarrow \left( \left\lceil \frac{1000}{33-1} \right\rceil \times 200 \right) + 1000$$

$$\Rightarrow ( \lceil 31.25 \rceil \times 200 ) + 1000$$

$$\Rightarrow 32 \times 200 + 1000$$

$$\Rightarrow \underline{7400} \text{ Block transfer.}$$

## 1.2 Merge Join

$$B_s \text{ (Sorting cost)} = b_r * \left( 2 \left\lceil \log_{m-1} \left( \frac{b_r}{m} \right) \right\rceil + 2 \right) + b_s * \left( 2 \left\lceil \log_{m-1} \left( \frac{b_s}{m} \right) \right\rceil + 2 \right)$$

$$b_r = 1000, \quad b_s = 200, \quad m = 33$$

Since data is not sorted  $\rightarrow$  need to calculate  $B_s$

$$B_s = 1000 * \left( 2 \left\lceil \log_{33-1} \left( \frac{1000}{33} \right) \right\rceil + 2 \right) + 200 * \left( 2 \left\lceil \log_{33-1} \left( \frac{200}{33} \right) \right\rceil + 2 \right)$$

$$B_s = 1000 * \left( 2 \left\lceil \log_{32} 30 \cdot 30 \right\rceil + 2 \right) + 200 * \left( 2 \left\lceil \log_{32} 6 \cdot 06 \right\rceil + 2 \right)$$

$$= 1000 * (2(1) + 2) + 200 * (2(1) + 2)$$

$$\Rightarrow 4000 + 800$$

$$B_s = \underline{\underline{4800}} \text{ writing cost Block transfer.}$$

$$\Rightarrow \text{Total cost} \Rightarrow B_s + b_r + b_s$$

$$\Rightarrow 4800 + 1000 + 200$$

$$\underline{\underline{\text{Cost}}} \Rightarrow \underline{\underline{6000}} \text{ Block transfer.}$$

### (1.3) Hash Join (Recursive Partition)

$$\text{Cost} \Rightarrow 2(b_r + b_s) \left\lceil \log_{m-1} (b_s) - 1 \right\rceil + b_r + b_s \Rightarrow \underline{\underline{b_s \leq b_r}}$$

$$\text{Since } 200 \leq 1000 \Rightarrow b_s \leq b_r.$$

$$\Rightarrow \text{Cost} \Rightarrow 2(1000 + 200) \left\lceil \log_{33-1} (200) - 1 \right\rceil + 1000 + 200$$

$$\Rightarrow 2(1200) \left\lceil \log_{32} 200 - 1 \right\rceil + 1200$$

$$\Rightarrow 2400 \left\lceil 1.5 - 1 \right\rceil + 1200$$

$$\text{Cost} \Rightarrow 2400(1) + 1200$$

$$\underline{\underline{\text{Cost}}} = \underline{\underline{3600}} \text{ Block transfer.}$$



1.4 Hash Join (No recursive partition)

$$\begin{aligned} \text{Cost} &= 3(br + bs) \\ &= 3(1000 + 200) \\ &= 3 \times 1200 \\ \text{Cost} &\Rightarrow 3600 \text{ block transfers.} \end{aligned}$$

1.5 If there are infinity of memory

$\Rightarrow$  Block Nested Loop Join :-  $\text{Cost} = br + bs = 1000 + 200 = 1200$   
 $\Rightarrow (M \geq br + bs) \Rightarrow \text{Seeks} = 2 \text{ seeks.}$

$\Rightarrow$  Hash Join  $\Rightarrow$  If entire table can be kept in main memory, no partitioning require

$\Rightarrow \text{Cost} \Rightarrow br + bs \Rightarrow 1000 + 200 \Rightarrow 1200 \text{ block transfers}$

$\vee$  Hash Join  $\Rightarrow$  No recursive pattern is independent of  $M$

$\Rightarrow$  Block Nested Loop Join is preferred.

\* Block Nested Loop Join is preferable, if relations fits entirely in memory. Cost will be  $br + bs$  i.e., 1200 blocks transferable.

Because: - we can use longer relation as outer relation and read the block. The buffer will hold entire smaller relation one page for reading longer one for outer

## 2. Equivalent Expression (5 points)

There are four relations as following:

instructor (id, name, dept\_name, salary),

department (dept\_name, building, budget),

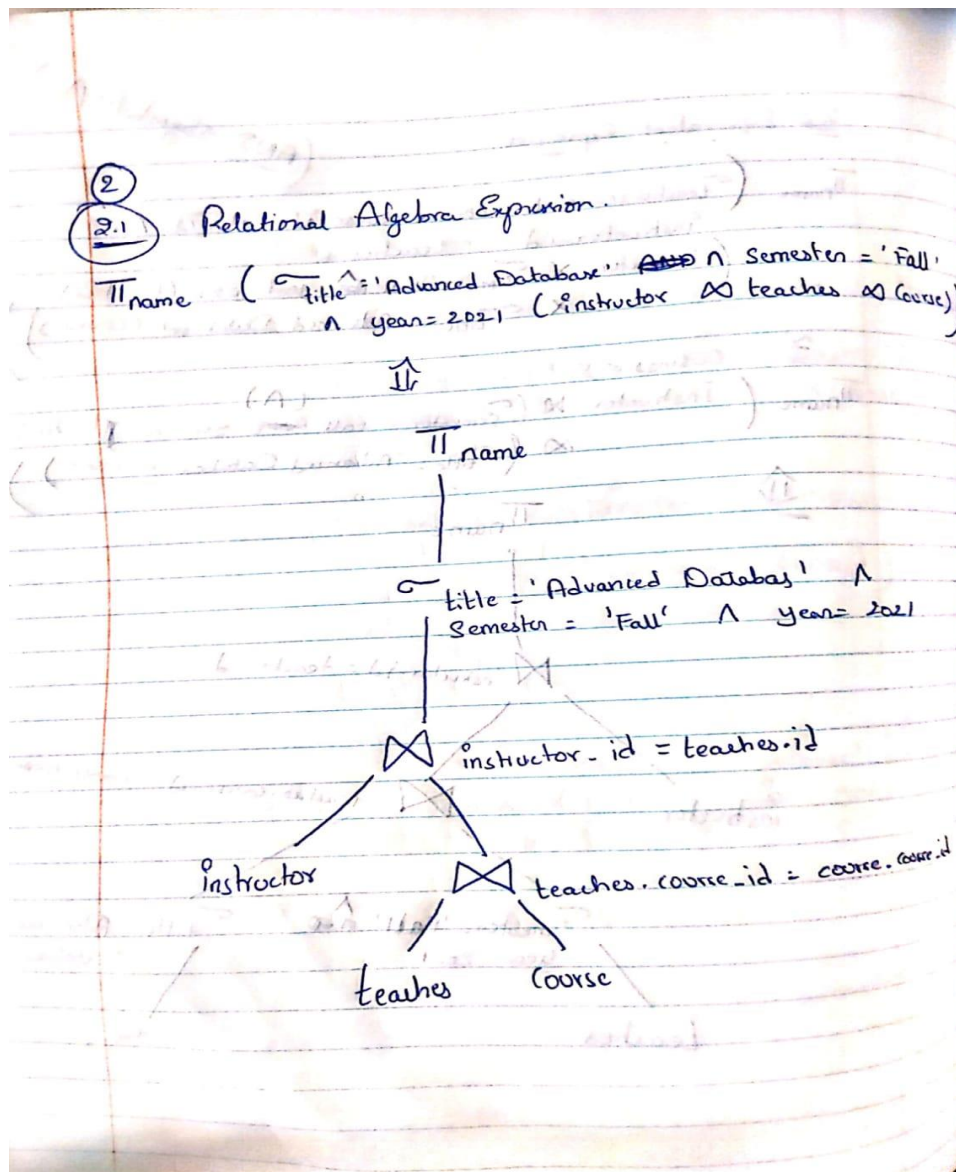
teaches (id, course\_id, sec\_id, semester, year),

course (course\_id, title, credits).

Query: Find the instructor's name who teaches 'Advanced Database' in Fall semester, 2021.

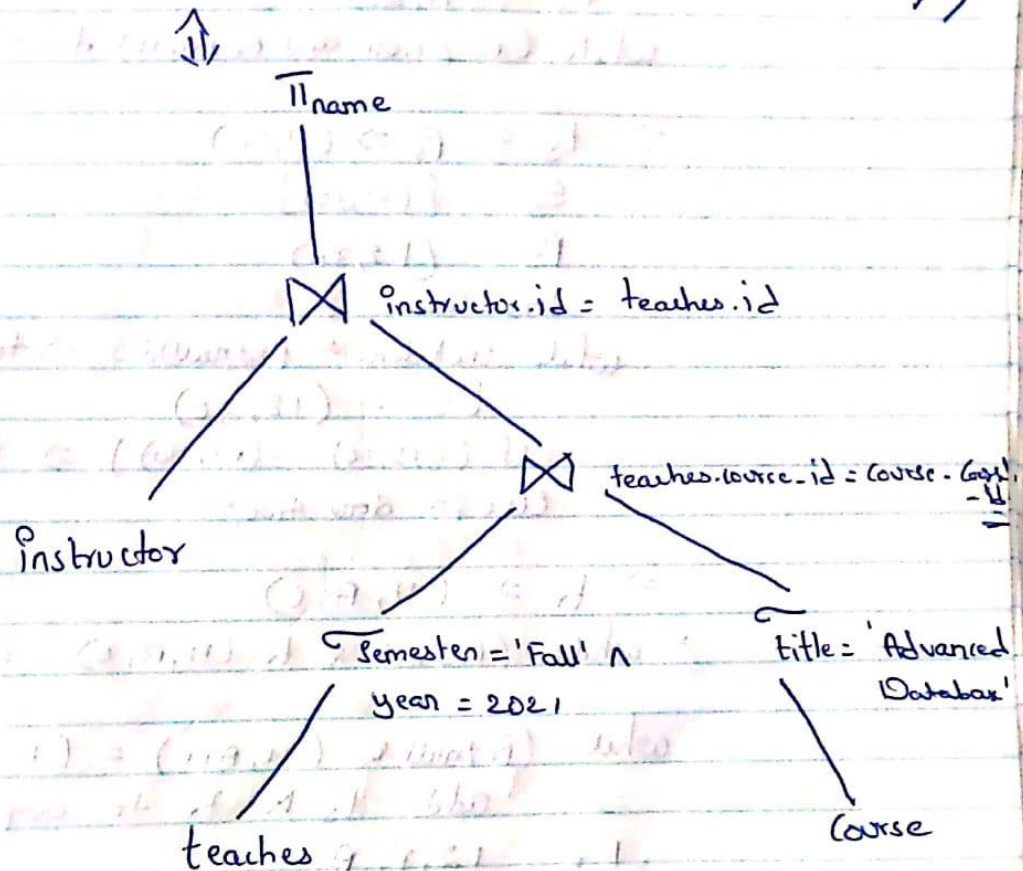
2.1 Find the relational algebra expression of this SQL command. (2 points)

2.2 According to 2.1, find the equivalent expression? And show how this equivalent expression is better than the expression on 2.1 (3 points)



## 2.2 Equivalent Expression

$$\pi_{\text{name}} \left( \text{instructor} \bowtie \left( \neg_{\text{Semester} = \text{'Fall'} \wedge \text{year} = 2021} (\text{teaches}) \right) \right. \\ \left. \bowtie \left( \neg_{\text{title} = \text{'Advanced Database'}} (\text{course}) \right) \right)$$



With Cartesian (X)

$$\pi_{\text{name}} \left( \neg_{\text{teaches.course.id} = \text{course.course.id} \wedge \text{instructor.id} = \text{teach.id}} \left( \text{instructor} \times \left( \neg_{\text{Semester} = \text{'Fall'} \wedge \text{Year} = 2021} (\text{teaches}) \right) \times \right. \right. \\ \left. \left. \neg_{\text{title} = \text{'Advanced Database'}} (\text{course}) \right) \right)$$



### 3. Merge Join Algorithm (10 points)

```
pr := address of first tuple of r;  
ps := address of first tuple of s;  
while (ps ≠ null and pr ≠ null) do  
  begin  
    ts := tuple to which ps points;  
    Ss := {ts};  
    set ps to point to next tuple of s;  
    done := false;  
    while (not done and ps ≠ null) do  
      begin  
        t's := tuple to which ps points;  
        if (t's[JoinAttrs] = ts[JoinAttrs])  
          then begin  
            Ss := Ss ∪ {t's};  
            set ps to point to next tuple of s;  
          end  
        else done := true;  
      end  
    tr := tuple to which pr points;  
    while (pr ≠ null and tr[JoinAttrs] < ts[JoinAttrs]) do  
      begin  
        set pr to point to next tuple of r;  
        tr := tuple to which pr points;  
      end  
    while (pr ≠ null and tr[JoinAttrs] = ts[JoinAttrs]) do  
      begin  
        for each ts in Ss do  
          begin  
            add ts ⋈ tr to result;  
          end  
        set pr to point to next tuple of r;  
        tr := tuple to which pr points;  
      end  
    end.  
  end.
```

---

Figure 15.7 Merge join.



Using Figure 15.7 Merge join with the following samples relation R and S

A1	A2	A3
11	A	C
12	F	A
12	L	K
14	T	P
15	I	O
16	P	L
17	K	C

R	A1	A4	S
	11	30	
	12	30	
	12	20	
	14	40	
	14	10	
	16	50	

How many rounds for the outer while loop?

Fill in the tuple on R that tR points to, tuple on S that tS points to and set SS after the end of each round.

Round#	tuple on R	tuple on S	Ss

3)

Round 1

$P_r = (11, A, C)$

$P_s = (11, 30)$

while ( $P_s \neq \text{null}$  and  $P_r \neq \text{null}$ ) do  $\Rightarrow$  true

$\Rightarrow t_s = P_s \Rightarrow (11, 30)$

$S_s = \{(11, 30)\}$

$P_s = (12, 30)$

done = false;

while (not done &  $P_s \neq \text{null}$ ) do  $\Rightarrow$  true.

$t_s = (12, 30)$

if  $((12, 30) = (11, 30)) \Rightarrow$  false

else  $\Rightarrow$  done = true;

$\Rightarrow t_r = (11, A, C)$

$\times$  while ( $P_r \neq \text{null}$  &  $t_r(11, A, C) < t_s(11, 30)$ )  $\Rightarrow$  false

while ( $P_r \neq \text{null}$  &  $(11, A, C) = (11, 30)$ )  $\Rightarrow$  true

$\Rightarrow$  add  $t_s$  &  $t_r$  to output.

$P_r = (12, F, A)$

$t_r = P_r = (12, F, A)$

end of Round 1

Round 2

while ( $P_r \neq \text{null}$  &  $P_s \neq \text{null}$ ) do  $\Rightarrow$  true

$t_s = (12, 30)$

$S_s = \{(12, 30)\}$

$P_s = (12, 20)$

done = false;

while (not done &  $P_s \neq \text{null}$ ) do  $\Rightarrow$  true

$t_s = (12, 20)$

if  $((12, 20) = (12, 30)) \Rightarrow$  true

$S_s = \{(12, 30)\} \cup \{(12, 20)\}$

$S_s = \{(12, 30) \& (12, 20)\}$

$P_s = (14, 40)$

$t_r = (12, F, A)$

$\times$  while ( $P_r \neq \text{null}$  &  $t_r(12, F, A) < t_s(12, 20)$ ) false

while ( $P_r \neq \text{null}$  &  $(12, F, A) = (12, 30)$ ) true

break Add  $t_s$  &  $t_r$  to output.

$P_r = (12, L, K)$

$t_r = (12, L, K)$

end of Round 2

Round 3

while ( $P_r \neq \text{null}$  &  $P_s \neq \text{null}$ ) true

$t_s = (14, 40)$

$S_s = \{(14, 40)\}$

$P_s = (14, 10)$

done = false

while (true)  $\Rightarrow$  true

$t_s = (14, 10)$

if  $((14, 10) = (14, 40))$  true

$S_s = S_s \cup \{t_s\}$

$S_s = \{(14, 40), (14, 10)\}$

$P_s = (16, 50)$

$t_r = (12, L, K)$

while ( $P_r \neq \text{null}$  &  $12 < 14$ )  $\Rightarrow$  True

$P_r = (14, T, P)$

$t_r = (14, T, P)$

while  $\Rightarrow$  false

while ( $P_r \neq \text{null}$  &  $14 = 14$ ) True

add  $t_s$  &  $t_r$  to output

$P_r = (15, I, O) = t_r$

end of Round 3

Round 4

$t_s = (16, 50)$

$S_s = \{(16, 50)\}$

$P_s = \text{Null}$

done = false

while ( $P_s \neq \text{null}$ )  $\Rightarrow$  false

$t_r = (15, I, O)$

while ( $P_r \neq \text{null}$  &  $15 < 16$ )  $\Rightarrow$  true

$\Rightarrow P_r = (16, P, L)$

$\Rightarrow t_r = (16, P, L)$

while (false)

while ( $P_r \neq \text{null}$  &  $16 = 16$ ) True

for  $\Rightarrow$  add  $t_s$  &  $t_r$  to output

$P_r = (17, K, C)$

$t_r = (17, K, C)$

end of Round 4

Round 5

$P_r = \text{null}$

while (false)

End

Round #	tuple on R	tuple on S	Ss
1	(12, F, A)	(12, 30)	{(11, 30)}
2	(12, L, K)	(14, 40)	{(12, 30), (12, 20)}
3	(15, I, O)	(16, 50)	{(14, 40), (14, 10)}
4	(17, K, C)	NULL	{(16, 50)}
-	-	-	-

A total of 4 Rounds for outer while loop.

Output

A <sub>1</sub>	A <sub>2</sub>	A <sub>3</sub>	A <sub>4</sub>
11	A	L	30
12	F	A	30
12	F	A	20
14	T	P	40
14	T	P	10
16	P	L	50

### Block Nested-Loop Join (~~r~~ $\Theta$ s)

The cost =  $\left(\left\lceil \frac{b_r}{M-1} \right\rceil * b_s\right) + b_r$  block transfers (disk accesses)

### Merge Join (~~r~~ $\Theta$ s)

$$B_s (\text{Sorting Cost}) = b_r * \left(2 \left\lceil \log_{M-1} \left(\frac{b_r}{M}\right) \right\rceil + 2\right) + b_s * \left(2 \left\lceil \log_{M-1} \left(\frac{b_s}{M}\right) \right\rceil + 2\right)$$

The cost =  $B_s + b_r + b_s$  block transfers (disk accesses)

### Hash Join (~~r~~ $\Theta$ s): No recursive partition

The cost =  $3(b_r + b_s)$  block transfers (disk accesses)

### Hash Join (~~r~~ $\Theta$ s): Recursive partition

The cost =  $2(b_r + b_s)[\log_{M-1}(b_s) - 1] + b_r + b_s$  block transfers (disk accesses); where  $b_s \leq b_r$